

The Mortality and Medical Costs of Air Pollution- Review

Author: Daniel Saggau — daniel.saggau@campus.lmu.de

Supervisor: Nadzeya Laurentsyevea

15/03/2021

1 Introduction

This research paper examines heterogeneity in mortality, health care and medical costs caused by pollution. Specifically, this paper focuses on heterogeneity in age, suggesting that a small subset of the elderly population is most vulnerable to pollution. This paper instruments for air pollution by leveraging changes in wind direction to estimate the causal impact of PM 2.5 on health using a subset of a sample of 40 billion observations. Using this instrument, they conduct a quasi-experimental investigation. Their research ensures robustness by looking at fixed effects, climate variables, changes in country's daily wind direction. The authors use various survival models, also looking at two machine learning methods namely a cox-lasso model and for reference a random survival forest. Individual level estimates are aggregate to get a country level measure. Their research estimates costs caused by exposure to pollution. The argument suggests that the 'mortality burden' is predominately placed within the elderly population with five to ten or two to five years remaining. Further, the authors introduce an estimate for life years left, arguing that life expectancy is a better indicator for vulnerability to pollution. This study disentangles variation in exposure using the generic machine learning approach by Chernozhukov et al. (2018). We can see that there is a near 0 effect among 75 percent of our sample looking at elderly beneficiaries. But the CDDF approach allows us to disentangle heterogeneity, and shows that from the 5 percentage threshold of the most vulnerable elderly population onward, pollution exposure has a very severe effect.

2 Data

This section looked at the different dataset. The authors looked at (1) air pollution, (2) atmospheric conditions and (3) mortality, morbidity and medical costs.

Air pollution is measured looking at PM 2.5 given in the Air Quality System Database, provided by the EPA starting in 1999. Other measures are also included in the dataset such as sulfur dioxide, ozone, nitrogen dioxide and carbon monoxide. For atmospheric conditions, we use information on wind speed and wind direction for the years 1999-2013 given within the North American Regional Reanalysis (NARR) daily reanalysis data. Information on Wind is reported at the level of a 32 by 32 kilometer grid with two vector pairs. The authors use interpolation between grid points to estimate the components entailing two different directions. Average measures are subsequently transformed into wind direction and speed at the county level. To supplement this information, the authors obtain information on temperature and precipitation data from Schlenker and Roberts (2009). Again, measures are averaged to get county-day level measures. Mortality, morbidity and medical costs are measured by using medicare administrative data. The sample focuses on beneficiaries between 65 and 100 years old. To calculate medical costs, the authors use the Medicare Provider Analysis and Review File (MedPAR). The unit of observation for these costs is the individual patient level. For the subsequent survival models, the authors also look at individual chronic conditions, to add to existing studies only looking at age and sex. Here, this paper ensures that there is no mortality displacement by using other time windows (5,14,28 days).

3 Methods (Empirical Strategies)

The methods and results are divided into three subsections, firstly looking at (1) mortality and health care use, the looking at (2) life years lost and thirdly (3) looking at treatment effect heterogeneity.

3.1 Mortality and health care use

Deryugina et al. suggest the following equation:

$$Y_{cdmy} = \beta \times PM2.5 + X'_{cdmy} \times \gamma + \alpha_c + \alpha_{as} + \alpha_{my} + \epsilon_{cdmy} \quad (1)$$

Here, Y is the outcome variable. As mentioned we are looking at mortality and health care use. For mortality we look at death per million beneficiaries as a three day total. For health care use the authors look at emergency room visits. The authors use estimates for ER visits that lead to hospital submission and as a placebo also looking at planned admissions. This measure should be independent of pollution. The study accounts for various fixed effects such as country level fixed effects, state by month level fixed effects and month by year level fixed effects. Fixed effects are notated as α_c for state by county, α_{sm} for state by month, and α_{my} for month by year fixed effects.

X is our specification for the remaining control variables. The authors use health care spending as a measure for medical costs and also look at health care. These measures are corrected for confounding factors. Subsequently, their study take the estimates and aggregate them to the country level, weighting by population per capita as the dependent variable.

Further the authors cluster air pollution monitors to ensure that pollution not dependent on the location of the monitor using k-means algorithm and creating 100 spatial groups. This approach allows for robustness against bias caused by selective placing of pollution monitors. As mentioned, the paper uses an IV-strategy using daily wind direction. The first stage equation is defined as: Geographic variation is enabled by using the interaction coefficient β .

$$PM2.5_{cdmy} = \sum_{g \in G} \sum_{b=0}^2 \beta_b^g 1[G_c = g] \times WINDDIR_{cdmy}^{90b} + X'_{cdmy} \sigma + \alpha_c + \alpha_s m + \alpha_m y + \epsilon_{cdmy} \quad (2)$$

The excluded variables are defined within the indicator function. $WINDDIR$ is equal to 1 if the direction in the country falls within the bandwidth of the 90 degree interval and 0 otherwise. The alpha measures are the defined in the same manner as for the first equation. Further, the authors argue that non-local pollution (traveling) will be the driving force of variation within the second equation because the non-local effects will be more likely to have

similar effects as opposed to the local effects which are more susceptible to differ. The effect of wind direction is confined by using bins. The specification for the remaining control variables is X , again. Further, this specification captures variation dependent on wind changes and their affect on local pollution.

3.2 Life-Years Lost

To account for mortality displacement outside of the mentioned window of days, this paper also estimates number of life years lost due to pollution exposure. Here, we want to disentangle the effect of pollution as opposed to average life years. To measure life years lost usually one needs information on counterfactual life expectancy which is unobserved. Previous studies use e.g. counterfactual life expectancy via population life tables. The solution proposed in this paper is using a rich set of different health and non-health characteristics to add to existing estimates only looking at sex and age. The authors also comparing propose estimates based on machine learning methods which allow us to enable feature selection. As mentioned the authors use a non-parametric model, a survival forest, and a regularization based model, a least absolute shrinkage and selection operator (LASSO) with a cox proportional hazard model. The regularization method induces a penalization term which is guarded by the parameter λ , weighting the trade off between bias and variance. The lasso model is a commonly used too due to its properties allowing one to induce further scarcity. The authors report all estimates for the different methods (Medicare FFS average, cox, cox with more covariates, survival random forest, cox lasso). There is a substantial drop in expected life expectancy when accounting for these chronic conditions and other non-health characteristics. The machine learning methods propose the lowest predicted life expectancy in years.

3.3 Treatment Effect Heterogeneity

To disentangle heterogeneity in our treatment this paper uses the CDDF approach proposed by Chernozhukov et al. (2018) to look at heterogeneity in age for the model looking at mortality. The authors focus on estimating the best linear predictor of conditional average treatment effect under general conditions. The authors use a gradient boosted decision tree

as a machine learning method for identification of treatment heterogeneity. We can get a proxy predictor based on this machine learning method. Using this proxy predictor one can now construct a weighted regression, generate a pooled sample with control and treatment group observations.

The introduce the following notation:

$$Died_{it} = \alpha + \beta_1(T_{ik} - \hat{p}(Z_{it}) + \beta_2(T_{ik} - \hat{p}(Z_{it}))(\hat{S}(Z_{it}) - \bar{\hat{S}}) + \theta \hat{Died}^C(Z_{it}) + \epsilon_{it} \quad (3)$$

Where $\hat{S}(Z_{it})$ is the proxy predictor for mortality. The outcome died is equal to 1 if the patient died (and 0 otherwise), $\hat{p}(Z_{it})$ is the propensity score of treatment. $\bar{\hat{S}}$ is the average of the proxy predictor. $\hat{Died}^C(Z_{it})$ are our control variables. Henceforth, the β_1 in the first part of the regression looks at the score of being in a treatment (T_{it}) versus predicted treatment($\hat{p}(Z_{it})$). The second part of the regression looks at the same relationship but multiplying this term with $\hat{S}(Z_{it}) - \bar{\hat{S}}$ where $\hat{S}(Z_{it})$ is the estimated deaths in treatment versus estimated deaths in control group and $\bar{\hat{S}}$ is average across the entire sample. As suggested by the authors, CCDF show that those terms are the best linear predictor of the conditional average treatment effect, allowing us to use the proxy predictor for this analysis. Next, the authors introduce a sorted group average treatment effect equation. The notation is as follows:

$$Died_{it} = \alpha + \sum_{k=1}^{\gamma} \gamma \times k(T_{ik} - \hat{p}(Z_{it})) * 1(G_k) + \theta \times \hat{Died}^C(Z_{it}) + \epsilon_{it} \quad (4)$$

Here, the paper uses an indicator function $1(G_k)$ see whether the prediction lies within within the kth interval or not. Of special interest is the gamma parameter, an estimate for the ATE for the group k. Two further challenges are addressed. On the one hand the authors down-sample because the probability of dying is small and fails poorly without doing so. The authors argue that computational constraints make it unable to include country, state-by-month and month by year fixed effects in this regression. Their research employs subgroups to estimate the equations and averages (using 250 subsets).

4 Results

There are 10 result tables that are included in the main paper. Due to time constraints, we won't go into detail for every table. For the results on mortality by age group the authors present both OLS and IV estimates, while one should note that OLS indicate significant bias. Same relationship holds for the health care usage results. The results of PM 2.5 exposure on estimates for the life years lost indicate that there is significant heterogeneity when accounting for more covariates. Using a 'back of the envelop calculation' the authors also propose a cost estimate. One should note that with the lower life years estimates, these costs also reduce (around \$76 billion dollars lower). The authors introduce the estimates for the different age groups proposed by the CDDF approach and subsequently also explore robustness by taking into account other pollution measures, weather controls and lags, too. In the conclusion, their study suggests that policy makers looking at the effect of exposure to pollution needs to account for vulnerable groups rather than focusing on improving the most polluted areas.

5 Conclusion

The contribution to the literature of this paper is suggested to be twofold. Firstly, this paper spends a sizable amount of defining a suitable estimate for life years left, partially by introducing machine learning for time to event studies for a economics study, looking beyond age and sex as determinants for life expectancy. Specifically, this paper uses two machine learning methods namely survival random forest and a cox-lasso model. Secondly, this paper uses this proposed generic machine learning for inference on heterogeneous treatment effects to a quasi-experimental study.

6 Review

This section will evaluate some of the most important or debatable components of this analysis. First, I will discuss the model evaluation for the predicted life expectancy. Afterwards, I will focus on the added value of using different machine learning methods. Subsequently, I will discuss other sources of heterogeneity that are ignored in this study. Lastly, I will take about

further measures to improve robustness and open questions.

6.1 Model Evaluation: Machine Learning for Survival Studies

Traditional model evaluation is omitted in this paper, justified by the argument that if there is no variation in the results, it does not matter how precise our estimates are. To ensure scientific rigor, one should have at least reported some sort of model performance for the machine learning methods. Further the argument that you don't need a concise estimate for identification is confusing. If we disregard precision and don't report any information on model evaluation, why do we build a machine learning model that is so marginally different from standard measures in the first place? Model evaluation would have been a substantial improvement to this paper because it would have allowed to assess the credibility of these different estimates allowing the reader to make his own judgment about the contribution of using these estimates apart from statistical significance of the subsequent regression results.

Model evaluation for classical classification tasks and survival tasks differentiate. To accommodate the right censored nature of our data (50 % of the individuals in our sample don't die within the time frame), we would use methods employing an estimation for our right censored observations e.g. via inverse weighted propensity estimates (IWPC). Not going into too much detail, we could have used the concordance statistics (in short c-statistics) or a brier score here to evaluate the performance. For more robustness, usually one would compare performance for the different thresholds within the data (typically: 25, 50, 75 percentage quantiles). Suppose the prediction would have been at either one of the extremes such an extremely high or low MSE, it would have been interesting to see and would have allowed some more insights into the relationship of different predictors.

6.2 Added Value of Survival Analysis via Machine Learning

The difference between the the survival random forest estimate and the cox-lasso is marginal (5.33 to 5.23). The difference between the cox-ph model and the cox-lasso model is slighter bigger (5.33 to 4.8) but still not a substantial jump as opposed to the jump between medicare FFS average and the cox model. Unfortunately, it is largely unclear whether these lower

predictions are actually a more accurate representation of survival, or whether these score are simply lower. It would have been really interesting to look at the different percentile thresholds for our data, given the subsequent arguments in the paper and the focus on vulnerability and heterogeneity within our sample. Further the argument that it only matters if the results expose differences is not very sound. Generally speaking, this does not imply that machine learning tools are not applicable here, solely that the contribution made here is marginal as compared to what existing methods propose. One should also note that there are various way to define statistical value of life. Examples include: Excess death, hazard fraction, relative risk, premature deaths, attributable deaths.(Hammitt et al., 2020). These various notions of statistical value in itself offer variety to create an estimate for life years remaining. It is not unlikely that there is also variation in these different estimates. Solely based on the research presented in this paper, one cannot exclude that the chosen measure was cherry picked to confirm a prior hypothesis or illustrate the relationship between traditional methods and machine learning estimates.

6.3 Contribution of the CDDF Approach and considerations

One central issue for the CDDF is that performance evaluation tools are missing. As mentioned above, the authors argue that due to the argument that prediction is not of central importance, one does not need to report performance evaluation metrics. These scores would have been interesting (even if not of central importance) to get an idea of actual predictive performance to put the entire analysis into perspective. This would have given the machine learning methods more attention and ground for justification. A single proxy estimator for treatment heterogeneity is hard to contextualize and evaluate. Rather than introducing a random forest for survival analysis, for the treatment effect heterogeneity one could have actually used these methods to see if there are any difference in generic machine learning methods. Especially also looking at e.g. bagging (Random Forest) and boosting methods (used here), it would have been insightful for completeness to see actual difference between machine learning methods even if so marginal. It would have been a more compelling addition as opposed to the survival random forest which was not in the center stage of the analysis. One should also note that one could should consider the difference between using generic

machine learning methods versus e.g. a causal random forest. One can use a causal random forests for lower dimensional data focusing on the conditional average treatment effects. Causal random forests are less prone to error caused by splitting uncertainty. Unfortunately, for this problem, the causal random forest is not applicable because we are working with a high dimensional dataset. One could argue that due to the sample size, one might be less vulnerable towards splitting error. Additionally, the CDDF approach shows that one can utilize generic machine learning methods without violating any assumptions, given the different trajectory. Irrespective, the uncertainty remains.

Nevertheless, one should also note that this approach is a substantial contribution. Not only do the results coalign with the results provided by the estimated life expectancy, but this method proposed novel opportunity to intertwine state of the art machine learning methods for economic research questions. Not many papers have actually utilized this novel approach, and this extensive research paper provides a further pathway of acceptance to use these statistical tools despite their shortcomings with respect to interpretation. Further, this also illustrates that there is room within economics for other machine learning tools concerned with similar trajectories. One example is interpretable machine learning, where the goal is to remove complexity from black box models and create models that are interpretable and usable for a broader audience, including policy researchers.

6.4 Heterogeneity in our Sample

Only use random 5 percent of FFS medicare sample, still amounting to 1.2 million individuals. This subset is justified in the paper as a tool for computational ease. While they do run the regression multiple times, it might have been nice to see other medicare beneficiaries. By disentangling the sample even further, one could potential discover even more heterogeneity. Accounting for factors like geography would have unfolded further variation in our treatment or at least influence our life expectancy estimates.

6.5 Other Minor Considerations

Feature Selection: Given that the paper uses so many control variables, it would have been nice to actually include a feature selection stage for the core regression. Here, the authors focus on feature selection for the main features and estimates but not for the other control variables. For transparency, this would have been a nice addition.

Autocorrelation: The problem with using lags is determining the lag structure of our data. Here, we simply use a single year lag which is somewhat naive, because we don't know whether the medicare data would actually translate in that manner. This is a potential source of miss-specification. Including an analysis on this issue in the appendix, would have allowed for more transparency.

Atmospheric science models: As suggested by the authors themselves, due to the vast number of observations and data, it is computationally not feasible to apply advanced tools from atmospheric science models here. Henceforth one can only speculate how these models would have performed.

7 Conclusion

In conclusion, this paper does spark an interesting discourse about pollution exposure, shifting the narrative from focusing on improving the most polluted areas to focusing on the most affected areas. Essentially, this study does provide systemic evidence for heterogeneity in treatment, suggesting that the most vulnerable groups within the elderly are more affected by pollution. Irrespective, there are various things that could have been done to enhance scientific rigor. The biggest shortcoming was the implementation of the various machine learning methods. There was no substantial improvement for the life expectancy estimates by using these tools as compared to the extensive cox-model for the estimated life expectancy estimates. One instrumental consideration is the disregard for actual model evaluation, reducing their credibility. The argument that precision is only important for identification when results differ is not very rigorous.

With respect to the second machine learning method for treatment effect heterogeneity one

should also consider the lack of scope in the treatment heterogeneity. We clearly see strong variation in age groups when it comes to treatment heterogeneity, but this is just one of many instrumental factors to consider. Nevertheless, one should also note that this research paper is a tremendous contribution to existing research. This study truly shows the versatility of machine learning tools for applications in a systematic manner. The CDDF approach illustrates great potential for future research to disentangle treatment heterogeneity for other factors such as in this case demographic and socio-economic variation.

Unanswered Questions

One of the many interesting questions that remain unanswered is the effect of pollution in other shares of the populations. Even looking at the new research paper published by the mostly same research team (apart from one coauthor), there is no consideration for other subgroups. This is a pivotal shortcoming, given the strong policy statements given in this and the subsequent paper by these authors. One interesting addition would be looking at infant mortality and the impact of pollution on younger age groups. Is it ethical to ground policy decisions solely based on analyses using subgroups of the population? In other terms, is it fair to make pollution reduction policies without accounting for other subgroups susceptible to pollution exposure? Further, there are more convoluted aspects that remain hard to disentangle. It remains uncertain whether we will be able to causally disentangle the impact of pollution accounting for true long-run heterogeneity within exposure. Ultimately, the question remains whether these insights will really be causal or merely a mere approximation of an causal effect, capturing some other sort of variation due to omitted factors.

8 References

- Chernozhukov, V., Demirer, M., Duflo, E., & Fernandez-Val, I. (2018). Generic machine learning inference on heterogeneous treatment effects in randomized experiments (No. w24678). National Bureau of Economic Research.
- Deryugina, T., Heutel, G., Miller, N. H., Molitor, D., & Reif, J. (2019). The mortality and medical costs of air pollution: Evidence from changes in wind direction. *American Economic*

Review, 109(12), 4178-4219.

Deryugina, T., Miller, N., Molitor, D., & Reif, J. (2021). Geographic and socioeconomic heterogeneity in the benefits of reducing air pollution in the United States. *Environmental and energy policy and the economy*, 2(1), 157-189.

Hammitt, J. K., Morfeld, P., Tuomisto, J. T., & Erren, T. C. (2020). Premature deaths, statistical lives, and years of life lost: identification, quantification, and valuation of mortality risks. *Risk Analysis*, 40(4), 674-695.